

Lecture 5:

Reinforcement Learning: Control

ECS7002P - Artificial Intelligence in Games

Diego Perez Liebana - diego.perez@qmul.ac.uk

Office: CS.301



<https://gaigresearch.github.io/>

Queen Mary University of London


Outline

Model Free Control in RL


Exploration versus Exploitation

Advanced Materials: $TD(\lambda)$ and SARSA

Key reminders




$$v_\pi(s) = \sum_{a \in \mathcal{A}} \pi(a|s) q_\pi(s, a)$$




$$q_\pi(s, a) = \mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a v_\pi(s')$$

Key reminders





$$v_*(s) = \max_a q_*(s, a)$$



$$q_*(s, a) = \mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a v_*(s')$$

Key Reminder - Policy Iteration



How to obtain π^* ? We iterate through the following two steps:

1. **Policy evaluation:** Estimate v_π
Dynamic Programming in Model-based
(e.g. Iterative policy evaluation)
2. **Policy improvement:** Generate $\pi' \geq \pi$
e.g. Greedy policy improvement

Q? Does Policy Iteration get stuck in Local Maxima? **No.**

Bellman expectation equation has v_π as a fixed point (contraction mapping theorem). <https://runzhe-yang.science/2017-10-04-contraction/>

Q? What else can we use to do *Policy Evaluation*?

Reinforcement Learning: Control

Model Free Control in RL


On vs Off-policy

On-policy: learn about policy π using π to sample

- On-policy Monte Carlo Control
- On-policy Temporal-Difference Learning (SARSA)


Off-policy: learn about policy π using π' to sample

- Off-policy learning (Q-Learning)



Model-Free Policy Iteration


1. **Policy evaluation:** Estimate v_π
Dynamic Programming (e.g. Iterative policy evaluation)
2. **Policy improvement:** Generate $\pi' \geq \pi$
e.g. Greedy policy improvement



We are going to play with what we use for Policy Evaluation and Improvement for the agent's behaviour. In **model-free**, can we do?

- Policy Evaluation: Monte-Carlo policy evaluation ($V = v_\pi$)
- Policy Improvement: Greedy policy improvement?

Model-Free Policy Iteration



A first problem: Monte-Carlo Policy Evaluation can't find the true value $V = v_\pi$ in model-free: we do not have full knowledge of $P_{ss'}^a$! Also, **improving** a policy (acting greedily with respect to V) requires the knowledge of the model ($P_{ss'}$):


$$\pi'(s) = \arg \max_{a \in A} R_s^a + P_{ss'}^a v(s')$$

The alternative is to use action-value function $Q(s, a)$

$$\pi'(s) = \arg \max_{a \in A} q_\pi(s, a)$$

Therefore, Monte Carlo can aim to approximate $Q_\pi(s, a)$. By caching a $Q_\pi(s, a)$ values (i.e. averaging returns), we can do control in a model-free setting, by picking the action that maximizes these q-values as policy.


On-policy Monte Carlo Control



- **Policy evaluation:** Monte Carlo policy evaluation, $Q = q_\pi$
- **Policy improvement:** Greedy policy improvement

Q? Is greedy policy the best policy to use?

On-policy Monte Carlo Control



- **Policy evaluation:** Monte Carlo policy evaluation, $Q = q_\pi$
- **Policy improvement:** Greedy policy improvement

Q? Is greedy policy the best policy to use?

No. By acting greedily, we do not explore the search space sufficiently enough.

Q? Why wasn't this a problem before with Dynamic Programming?

ϵ -Greedy Exploration


ϵ -Greedy Exploration:

- Simplest idea for ensuring continual exploration
- Ensures all actions are tried with > 0 probability
- With $1 - \epsilon$ probability, choose the greedy action.
- With ϵ probability, choose another action at random.

$$\pi(a \mid s) = \begin{cases} (1 - \epsilon) & a^* = \arg \max_{a \in A} q(s, a) \\ \epsilon & \text{otherwise.} \end{cases}$$

Monte Carlo Control

- **Policy evaluation:** Monte Carlo policy evaluation, $Q = q_\pi$
- **Policy improvement:** ϵ -Greedy policy improvement
- **Every Episode:** Perform Policy Improvement after every single episode: Collect all the steps during an episode, updating the q-values for the pairs (s,a) visited **only**, and improve the policy straight away.



Greedy in the Limit with Infinite Exploration (GLIE)

At some point, we want to stop exploring and pick the action that maximizes $Q(s,a)$ all the time! We need to find π^* .

A way to do this is to decrease the value of ϵ after each episode, until it reaches 0 (for example, $\epsilon \leftarrow \frac{1}{k}$).

```
1: procedure GLIE_MC
2:   Initialize  $q(s,a)$  arbitrarily,  $q(\text{terminal state}) = 0$                                 ▷ i.e.  $Q(s, a) = 0 \forall s \in S, a \in A$ 
3:   for all  $k \in (1 : N)$  do                                                                ▷ During  $N$  iterations of GLIE MC
4:     Generate an episode using  $\epsilon$ -greedy policy  $\pi$  ( $EP_\pi$ )
5:     for all  $s, a \in EP_\pi$  do
6:        $N(s_t, a_t) \leftarrow N(s_t) + 1$                                                     ▷ Increment visit counter
7:        $Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \frac{1}{N(s_t, a_t)} (G_t - Q(s_t, a_t))$         ▷ Update the Q-value of this pair  $(s_t, a_t)$ 
8:     end for
9:      $\epsilon \leftarrow \frac{1}{k}$ 
10:  end for
11: end procedure
```

GLIE Monte Carlo Control converges to the optimal action-value function:

$$Q(s, a) \rightarrow q^*(s, a)$$

Note: the algorithm improves $Q(s, a)$, which is used by the policy: policy improvement after every episode.

Off-policy Learning

Off-policy learning uses two different policies:

- Target policy ($\pi(a | s)$): policy that is being evaluated and improved.
- Behaviour policy ($\mu(a | s)$) used to sample the MDP, generating the sequence of $\{S_1, A_1, R_1, \dots, S_T\}$.

Why?

Off-policy Learning

Off-policy learning uses two different policies:

- Target policy ($\pi(a | s)$): policy that is being evaluated and improved.
- Behaviour policy ($\mu(a | s)$) used to sample the MDP, generating the sequence of $\{S_1, A_1, R_1, \dots, S_T\}$.

Why?

- Learn from other agents, even from humans.
- Re-use old policies used in the past ($\pi_1, \pi_2, \dots, \pi_{t-1}$).
- Learn about an *optimal* policy while following an *exploration* policy.
- Learn about *multiple* policies while following *one* policy.

Q-Learning


Q-Learning is probably the most famous Off-policy Learning in RL:

- Target policy ($\pi(a | s)$) is the greedy policy:

$$\pi(s_{t+1}) = \arg \max_{a' \in A} Q(s_{t+1}, a')$$

- Behaviour policy ($\mu(a | s)$) is the ϵ -greedy policy.
- Both policies improve on each iteration of the algorithm.
- We are learning action-values ($Q(s, a)$):

$$Q(s, a) \leftarrow Q(s, a) + \alpha(R + \gamma \max_{a'} Q(s', a') - Q(s, a))$$



We are trying to learn how to act optimally (target policy) while exploring (using the behaviour policy).

Q-Learning

```
1: procedure QLEARNING
2:   Initialize  $q(s,a)$  arbitrarily,  $q(\text{terminal state}) = 0$            ▷ i.e.  $Q(s, a) = 0 \forall s \in S, a \in A$ 
3:   for all  $k \in (1 : N)$  do                                           ▷ During  $N$  Episodes of QLearning
4:     for all  $s \in EP_\pi$  do                                           ▷ For all states in the episode
5:        $s' \leftarrow$  Choose an action  $a$  using the  $\epsilon$ -Greedy policy,  $\pi(s)$  (derived from  $Q(S, A)$ ).
6:       Determine the target to learn from with the max q-value:  $R + \gamma \max_{a'} Q(s', a')$ 
7:       Update  $Q(s, a)$ :  $Q(s, a) \leftarrow Q(s, a) + \alpha(R + \gamma \max_{a'} Q(s', a') - Q(s, a))$ 
8:        $s \leftarrow s'$ 
9:     end for
10:    Until  $s$  is terminal
11:  end for
12: end procedure
```

Q-Learning converges to the optimal action-value function:

$$Q(s, a) \rightarrow q^*(s, a)$$

Reinforcement Learning: Control

Exploration versus Exploitation

Exploration vs. Exploitation Dilemma

As we have seen in this and previous lectures, selecting actions involves a fundamental choice:


- *Exploitation*: Make the best decision based on current information.
- *Exploration*: Gather more information about the environment. This is: not choosing the best action found so far.

The objective is to gather enough information to make the best overall decision. The best long-term strategy may involve short-term sub-optimal selections.

There are different ways to explore:

- Random exploration: ϵ -greedy, Softmax, . . .
- Optimism in the face of uncertainty: estimate the uncertainty of a value, and prefer to explore those with higher uncertainty.

The Multi-Armed Bandit Problem



- A multi-armed bandit is a tuple $\langle A, R \rangle$.
- A is a known set of actions (arms).
- Set of unknown distributions $\{R_1, R_2, \dots, R_k\}$ of rewards, one per action.
- Played iteratively, during H action selections.
- Mean values of these reward distributions: $\{\mu_1, \mu_2, \dots, \mu_k\}$
- The goal is to maximize the sum of rewards (minimizing the loss).
- Each action is pulling one lever. How do you choose?

Regret

Regret is the opportunity loss (total, or for one step). How much did I lose because I did not choose the optimal action?

- Given the action value $Q(a)$ and the optimal value V^*

$$Q(a) = \mathbb{E}[r \mid a] \qquad V^* = Q(a^*) = \max_{a \in A} Q(a)$$

- The **regret** is the opportunity loss for one step:

$$l_t = \mathbb{E}[V^* - Q(a)]$$

- The **total regret** is the total opportunity loss:

$$L_t = \mathbb{E}\left[\sum_{t=1}^T (V^* - Q(a))\right]$$

- The objective is to minimize the total regret, which maximizes the cumulative reward.

Regret: gaps and counts


How do we count regret?

- The *count* $N_t(a)$ is the expected number of selections of action a .
- The *gap* Δ_a is $V^* - Q(a)$: difference in value between picking a and a^* .
- Total regret can be expressed as a function of *gaps* and *counts*:

$$\begin{aligned} L_t &= \mathbb{E}\left[\sum_{t=1}^T (V^* - Q(a))\right] \\ &= \sum_{a \in A} \mathbb{E}[N_t(a)](V^* - Q(a)) \\ &= \sum_{a \in A} \mathbb{E}[N_t(a)]\Delta_a \end{aligned}$$

- Therefore, a good algorithm produces small *counts* for large *gaps*, and viceversa, in order to minimize the total regret (L_t).
- **Q?** What's the problem?


Linear and Sublinear Regret



- Greedy: the algorithm never explores, the total regret is linear.
- Greedy with optimistic initialization
 - Initialize all $Q(a)$ to the maximum possible reward, then act greedily.
 - Still greedy, total regret is linear.
- (Constant)- ϵ greedy. It never stops exploring, hence the total regret is linear.
- (Decaying)- ϵ greedy: reduces slowly the value of ϵ at each step, it achieves sublinear regret.

Logarithmic regret

Can we do better?



- (Decaying)- ϵ greedy can achieve logarithmic regret ... **if we know the gaps in advance**, with the following decaying schedule:


$$c > 0 \quad d = \min_{a|\Delta_a > 0} \Delta_i$$

$$\epsilon_t = \min\left\{1, \frac{c |A|}{d^2 t}\right\}$$

- Logarithmic regret is actually the best we can do!
- **Goal:** an algorithm with logarithmic regret without knowing the gaps (Δ).

Bounds

Optimism in the face of uncertainty.




- Which action should we pick? The more uncertain we are about an state-action value, the more important is to explore that action.
- After picking an action, we are less uncertain about it, and more likely to pick another action.
- We keep until we build confidence on the action-value of each action.
- We know how to calculate the action-value, but how do we build the confidence?

Building the Bounds

How do we build this confidence?

- Estimate an upper confidence $U_t(a)$ for each action value that depends on the number of times a has been selected ($N(a)$).
 - Small $N(a)$: large $U_t(a)$ that implies uncertainty.
 - Large $N(a)$: small $U_t(a)$ that implies more accuracy.
- The action must be selected maximizing the Upper Confidence Bound (UCB):

$$a_t = \arg \max_{a \in A} \{Q_t(a) + U_t(a)\}$$



Upper Confidence Bounds

Theorem (Hoeffding's Inequality)

Let X_1, X_2, \dots, X_t be identically and independently distributed random variables in $[0, 1]$, and let $\bar{X}_t = \frac{1}{t} \sum_{\tau=1}^t X_\tau$ be the sample mean. Then:

$$\mathbb{P}[\bar{X}_t > \mathbb{E}[X] + u] \leq e^{-2tu^2}$$

This means: “What is the probability that the difference between the empirical and the actual mean is greater than u ?”. Or, in other words, “What is the probability of making a mistake greater than u when estimating the mean?”.

This theorem says that this probability is no more than e^{-2tu^2} , for **any distribution**, if the random variables are bounded in $[0, 1]$. In our case:

$$\mathbb{P}[Q(a) > Q_t(a) + U_t(a)] \leq e^{-2N_t(a)U_t(a)^2}$$

Deriving from this, we obtain that, when $t \rightarrow \infty$: $U_t(a) = \sqrt{\frac{2\log(t)}{N_t(a)}}$

UCB1

This leads to the UCB1 algorithm:

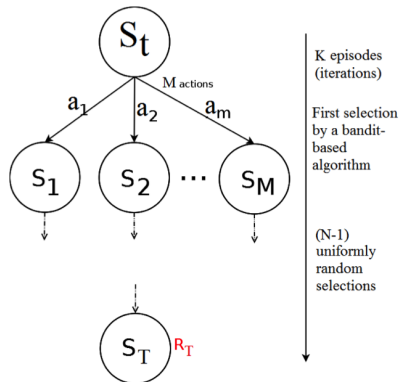
$$a_t = \arg \max_{a \in A} Q(a) + \sqrt{\frac{2 \log(t)}{N_t(a)}} = \arg \max_{a \in A(s)} Q(s, a) + C \sqrt{\frac{\ln N(s)}{N(s, a)}}$$

- $Q(s, a)$: Action-state value of action a from state s .
- $N(s)$: Times the state s has been visited.
- $N(s, a)$: Times the action a has been selected from state s .
- C : balances between *exploitation* and *exploration*:
 - Value of C is application dependent.
 - Example: single player games with rewards in $[0, 1]$: $C = \sqrt{2}$.
- UCB1 achieves **logarithmic** total regret.
- We don't need to know the gaps.
- There are many UCB variants, UCB1 is just one of them.
- Other theorems derive other UCB policies.

Flat UCB

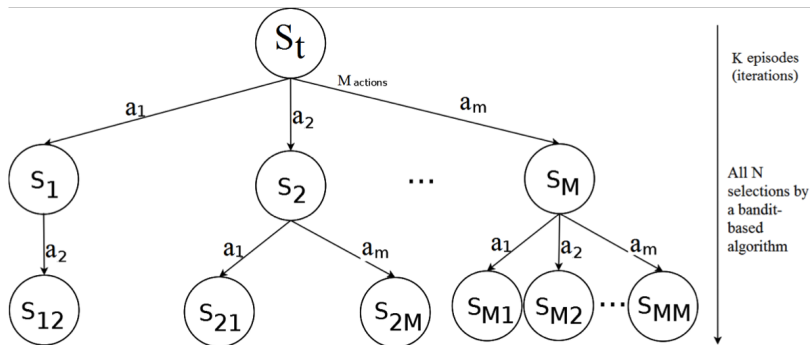
We can use UCB1 (or any other UCB policy) for searching the action-state space. For example, the *Flat UCB* algorithm:

- Iteratively, apply K episodes. For each one of them:
- Select the first action from S_t with UCB.
- Pick actions uniformly at random until reaching a terminal state (**roll-out**).
- This estimates state-action values $Q(s, a)$ from the state S_t .
- Note that the UCB policy improves at each episode.



Building a tree

By applying a UCB policy, we can add a node (that represents a state) at each iteration:



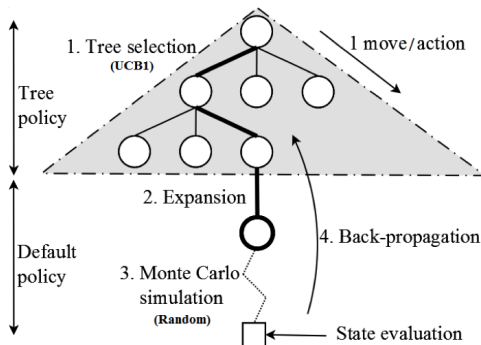
- The tree grows asymmetrically, towards the most promising parts of the search space.
- However, this is limited by how far can we look ahead into the future.
- If we add a node for each state visited during the random roll-outs, the tree would be too big!

Monte Carlo Tree Search (MCTS)

Monte Carlo Tree Search: adding Monte Carlo simulations after a new node is added to the tree.

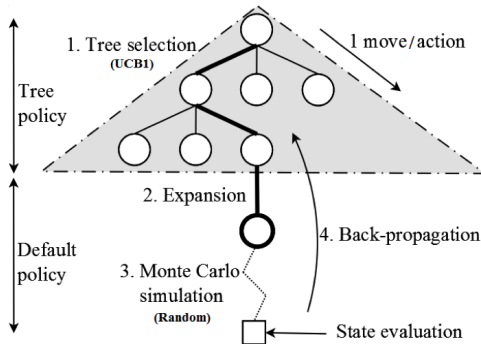
- 2 different policies are used on each episode:
 - **Tree policy** improves on each iteration. It is used while the simulation is in the tree. Some naming conventions:
 - UCT Algorithm: MCTS with any UCB tree selection policy.
 - Plain UCT Algorithm: MCTS with UCB1 as tree selection policy.
 - **Default policy** is fixed through all iterations. It is used while the simulation is outside the tree. Picks actions uniformly at random.
- On each iteration:
 - $Q(s, a)$ on each node in the tree is updated.
 - $N(s)$ and $N(s, a)$ on each node of the tree are updated.
 - Tree policy is based on $Q(s, a)$, thus it improves on each iteration.
- MCTS converges to the optimal search tree.

Monte Carlo Tree Search (MCTS)



1. **Tree Selection:** Following the tree policy (i.e. UCB1), navigate the tree until reaching a node with at least one child state not in the tree (this is, not all actions have been picked from that state in the tree).
2. **Expansion:** Add a new node in the tree, as a child of the node reached in the tree selection step.

Monte Carlo Tree Search (MCTS)




- 3. Monte Carlo simulation:** Following the default policy (picking actions uniformly at random), advance the state until a terminal state (or a pre-defined maximum number of steps). The state at the end of the simulation is evaluated (obtain the reward R).
- 4. Backpropagation:** Update the values of $Q(s,a)$, $N(s)$ and $N(s,a)$ of the nodes visited in the tree during steps 1 and 2.

Use Case: MCTS and the Game of Go



- 2500 years old 2-player game.
- Considered the hardest classic board game, and a challenge task for AI.
- Traditional game-tree search (minimax, alpha-beta search) has failed in Go: they can't reach human-play level.

Use Case: MCTS and the Game of Go



Use Case: MCTS and the Game of Go

Why is Go so difficult?

- The game is long (average of 200 moves; Chess: 40-50).
- Large branching factor (average of 250 legal plays/move; Chess: 35-40).
- But, primarily, lack of a good state evaluation function. It is not easy to evaluate how good or bad an intermediate state is:
 - A piece placed early in the game may have a strong influence later in the game, even if it will eventually be captured.
 - It can be impossible to determine if a group will be captured without considering the rest of the board.
 - Most positions are dynamic (there are always unsafe stones in the board).

How to approach it?

- Domain knowledge: find patterns in the board that represent strong plays.
- Use MCTS enhancements: AMAF, RAVE.
- Use delayed rewards.

- From 2016: MCTS + Deep Neural Networks (Alpha Go).

Use Case: MCTS and the Game of Go


MCTS was the first algorithm to achieve human-play level in the version with the small board.

1990	Abramson demonstrates that Monte Carlo simulations can be used to evaluate value of state [1].
1993	Brügmann [31] applies Monte Carlo methods to the field of computer Go.
1998	Ginsberg's GIB program competes with expert Bridge players.
1998	MAVEN defeats the world scrabble champion [199].
2002	Auer et al. [13] propose UCB1 for multi-armed bandit, laying the theoretical foundation for UCT.
2006	Coulom [70] describes Monte Carlo evaluations for tree-based search, coining the term Monte Carlo tree search.
2006	Kocsis and Szepesvari [119] associate UCB with tree-based search to give the UCT algorithm.
2006	Gelly et al. [96] apply UCT to computer Go with remarkable success, with their program MOGO.
2006	Chaslot et al. describe MCTS as a broader framework for game AI [52] and general domains [54].
2007	CADIAPLAYER becomes world champion General Game Player [83].
2008	MOGO achieves <i>dan</i> (master) level at 9×9 Go [128].
2009	FUEGO beats top human professional at 9×9 Go [81].
2009	MOHEX becomes world champion Hex player [7].

Timeline of events leading to the widespread popularity of MCTS.

Use Case: MCTS and the Game of Go

Popular MCTS Go Players:



Use Case: MCTS and the Game of Go

Go, MCTS and Deep Learning:



Acknowledgements

Additional materials: most of the materials for this course are based on the following resources.

- Prof. David Silver's course on Reinforcement Learning:
<http://www0.cs.ucl.ac.uk/staff/d.silver/web/Teaching.html>
- *Reinforcement Learning: An Introduction*, by Andrew Barto and Richard S. Sutton (2017 Edition):
<http://incompleteideas.net/book/bookdraft2017nov5.pdf>

All labs from now on: Monday 4-6pm, ITL ground floor.

- That includes the MCQ tests!!

MSc Projects.


- All projects have been designed aiming to response research questions that can be publishable in AI/Games conferences.
- Forward Models for Statistical Forward Planning methods:
 - Learning FMs, Abstract FMs, Incorrect FMs.
- Implementing and testing AI methods in a table-top board games:
 - SFP for wargames, automatic AI scripting, competition and cooperation in table-top board games.
- Designing and implementing a game description language:
 - Table-top board games, tile-based games, VGDL 3.0.

Reinforcement Learning: Control

Advanced Materials: $TD(\lambda)$ and SARSA

n-step Prediction

(the rest of this slide deck will not feature in the progress/MCQ test)



$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \cdots + \gamma^{T-t-1} R_T$$

n-step Return

The n-step returns look like this:

$n = 1$ (TD)	$G_t^{(1)} = R_{t+1} + \gamma V(S_{t+1})$
$n = 2$	$G_t^{(2)} = R_{t+1} + \gamma R_{t+2} + \gamma^2 V(S_{t+2})$
$n = 3$	$G_t^{(3)} = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \gamma^3 V(S_{t+3})$
\dots	\dots
$n = \infty$	$G_t^{(\infty)} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{T-1} R_T$

Therefore, we can define:

Definition (n-step Return)

$$G_t^{(n)} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n V(S_{t+n})$$

And:

Definition (n-step temporal-difference learning)

$$V(S_t) \leftarrow V(S_t) + \alpha(G_t^{(n)} - V(S_t))$$

Averaging n-step returns

Backups can be done not just toward any n-step return, but also toward any average of n-step returns. For instance, an average of a 2-step and 4-step return would be:

$$\frac{1}{2}G^{(2)} + \frac{1}{2}G^{(4)} = \frac{1}{2}G_t^{t+2}(V_t(S_{t+2})) + \frac{1}{2}G_t^{t+4}(V_t(S_{t+4}))$$

Q? Can we combine information from all time-steps?

TD(λ): averages n-step backups, weighting each one of them proportionally to λ^{n-1} ($\lambda \in [0, 1]$) and normalized by $(1 - \lambda)$ so the sum of weights = 1.

Definition (λ -Return)


$$(1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} G_t^n$$

λ -return


$$(1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} G_t^n$$

$\lambda = 0$: the overall backup reduces to the **first** component \rightarrow TD(0)

$\lambda = 1$: the overall backup reduces to the **last** component \rightarrow MC



λ -return




λ -Return: $G_t^\lambda = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} G_t^{(n)}$ (combines all n -step returns $G_t^{(n)}$)


Forward View TD(λ): $V(S_t) \leftarrow V(S_t) + \alpha(G_t^\lambda - V(S_t))$

Forward vs. Backward view

Forward view TD(λ): This is the TD view we have seen so far. It suffers from the same problems as MC (it's computed from complete episodes).




Backward view TD(λ) provides a way to do this with incomplete sequences. We need **eligibility traces**.



Backward view TD(λ)

Eligibility traces combine credit to most frequent and most recent states simultaneously:



Definition (Eligibility traces $E(s)$)

$$E_0(s) = 0$$

$$E_t(s) = \gamma\lambda E_{t-1}(s) + \mathbf{1}(S_t = s)$$

We keep an eligibility trace for every state s .

- Every time a state is visited, the eligibility trace increases.
- Eligibility traces decrease continuously. If a state receives no visits, it will decay up to a minimum.
- λ determines how rapidly the trace decays.

TD(λ)

TD(λ) updates value $v(s)$ for every state s , in proportion to TD-error δ_t and the eligibility trace $E_t(s)$.

$$\delta_t = R_{t+1} + \gamma v(S_{t+1}) - v(s_t);$$

$$E_t(s) = \gamma \lambda E_{t-1}(s) + \mathbf{1}(S_t = s)$$

$$v(s) \leftarrow v(s) + \alpha \delta_t E_t(s)$$

If $\lambda = 0$, $E_t(s) \leftarrow \mathbf{1}(S_t = s)$. This is an automatic, abrupt decay. Thus, we only care about the current state, so $v(s) \leftarrow v(s) + \alpha \delta_t$, which is equivalent to TD(0).

If $\lambda = 1$, there is no decay, we care about all states visited. This is (roughly) equivalent to every-visit MC.


If $0 < \lambda < 1$, less credit is given to δ_t errors from the past. The closer λ to 0, the more abrupt the decay of $E_t(s)$ becomes, and past δ_t errors have less effect on the update of $v(s)$.

SARSA


Same concept as in GLIE-MC, but using TD instead of MC.

This removes the need from simulating until the end of the episode.

$$Q(S, A) \leftarrow Q(S, A) + \alpha(R + \gamma Q(S', A') - Q(S, A))$$



SARSA



- **Policy evaluation:** Sarsa, $Q = q_\pi$
- **Policy improvement:** ϵ -Greedy policy improvement
- **Every Episode:** Perform Policy Improvement after every single episode.

SARSA

```
1: procedure SARSA
2:   Initialize  $q(s,a)$  arbitrarily,  $q(\text{terminal state}) = 0$ 
3:   for all  $k \in (1 : N)$  do
4:     for all  $s \in S$  do
5:       Choose  $a$  from  $s$  using  $\epsilon$ -Greedy policy,  $\pi(s)$ .
6:       Take action  $a$ , observe  $R$  and  $s'$ .
7:       Choose  $a'$  from  $s'$  using  $\epsilon$ -Greedy policy,  $\pi(s)$ .
8:        $Q(S, A) \leftarrow Q(S, A) + \alpha(R + \gamma Q(S', A') - Q(S, A))$ 
9:     end for
10:  end for
11: end procedure
```

▷ i.e. $Q(s, a) = 0 \forall s \in S, a \in A$
▷ During N iterations of SARSA
▷ For all states

SARSA converges to the optimal action-value function:

$$Q(s, a) \rightarrow q^*(s, a)$$

n-step SARSA, $\text{SARSA}(\lambda)$, $\dots \rightarrow$ same degree of control between MC Control and SARSA by tuning λ .