

Machine Learning Lecture 11

Features and dimensionality reduction

Dr. Ioannis Patras

Slides thanks: Tim Hospedales




1

Overview

- Feature Design
- Why Reduce Dimensions?
- Feature Selection and Methods
 - Filtering
 - Wrapper
 - Built-in
- Dimensionality Reduction
 - PCA
- Summary

2

Overview

- **Feature Design**  How to choose/design your features
- Why Reduce Dimensions?
- Feature Selection and Methods
 - Filtering  How to prune unhelpful features
 - Wrapper
 - Built-in
- Dimensionality Reduction  How to compress columns
 - PCA
- Summary

3

What do we mean by 'Features'

- Input database columns, aka Attributes, Dimensions.
- What data to include as the input to your learning procedure
 - Sometimes choice is out of your scope or obvious from domain/business setting
 - Sometimes choice is an opportunity for engineering / intuition
 - (Given that more attributes potentially provide more information, but potentially increase overfitting, memory and slow the computation)

4

Feature Choice

- E.g., You work at Zoopla, a real-estate web-site.
 - Your job is to add a feature that predicts prices for listed houses
 - What should you include?
 - Rooms, m² Area, Postcode, Distance to Transport/School, Crime...?
 - Floor type? Date of construction? Room shape? Wallpaper color...? Number of internal doors?
- Using linear/non-linear regression, you will find that every feature added increases R^2 / decreases RMSE.
 - How to decide when to stop?
 - Continuous increase of fit with irrelevant features due to over fitting
 - An option is to use cross-validation to decide if which features

5

Feature Engineering

- Input and Output for Machine Learning Algorithm
 - $\{\mathbf{x}_i=[x_1, \dots, x_d], y_i\}$
- How to convert real life data:
 - Into $\mathbf{x}_i=[x_1, \dots, x_d], y_i$?
 - What to store in your database?
- Domain Specific
 - Human Expertise



6

Feature Engineering

- It may make sense for your input to be some **transformation** of the raw data
 - Because your data may not be fixed length
 - Because the right non-linear transformation can make learning easier
 - Because the right low-dimensional transform could help avoid over-fitting.
- Designing this transformation:
 - **Feature engineering**



“Danger”

7

Feature Engineering: Examples

- E.g., Audio or accelerometer data
 - Full data is waveform
 - Probably don't want to use directly: Variable and high dimensions (E.g., 5sec/44KHz/stereo: Half million columns.)
- Features:
 - High-amplitude count (1d) ... Loud noise detector.
 - Zero-crossing count(1d) ... Activity recognition.
 - Fourier Coefficients (e.g., 128d) ... Music / Speech recognition
- **Case Study @ EECS:**
 - Activity + identity recognition on mobile phone accelerometer



8

Feature Engineering: Examples


- E.g., Images.
 - Full data is all the pixels
- Features:
 - Average brightness (1d) e.g., Day vs Night.
 - Color histogram (3 – 1000dims) e.g., Recognize objects
 - Histogram of Gradients (~128dim) e.g., Detect objects
 - Raw Pixels (10000dim+) e.g., Face recognition



9

Feature Engineering: Examples


- E.g., Internet Advertising
 - Full data is everything on your facebook profile.
 - Should we show an add about premium baby-items?
 - Engineer a feature aggregating high income related likes+posts & a feature aggregating baby/mother related likes+posts
- Could result in a simpler & more efficient model than if you threw everything in



10

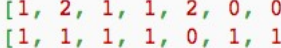
Feature Engineering: Text

- Text is a heavily studied/mined domain
 - Naturally: High dimensional & variable length.
 - How to represent as fixed length and low(-ish) dimensional?
- One option is meta-features
 - Feature Engineering of Document:
 - Length, average word length, punctuation frequency, (in)correct spelling ratio, emoticon density, uppercase/lowercase ratio, etc.
 - (6 dims in this example)
- Case Study: Using meta-features:
 - IM/SMS authors can be recognised



11

Feature Engineering: Text

- The most common choice for text is “Bag of Words”
- E.g., Raw data:
 - “John likes to watch movies. Mary likes movies too.”
 - John: 1, likes: 2, to: 1, ...
 - “John also likes to watch football games.”
 - John: 1, also: 1, likes: 1, ...
- Result is a “bag of words” vector \mathbf{f}

 - Length: # of words in the dictionary.
 - Row: Frequency of each word in a document
 - Sum of a row: Number of words in corresponding document
- Can also use bi-grams, trigrams, n-grams

```
[1, 2, 1, 1, 2, 0, 0, 0, 1, 1]
[1, 1, 1, 1, 0, 1, 1, 1, 0, 0]
```

12

Case Studies

EECS Work 😊

- Using Bag of Words + Linear Classifiers, previous project students did....
 - Predict bill passage from bill text.
 - Predict politician identity from speeches.
 - Predict movie genre from script/subtitles
 - Predict TV show from script/subtitles
 - Predict actor/director/scriptwriter from script/subtitles
 - Predict amazon product review rating + helpfulness
 - Predict sales rank/price from product description



☆☆☆☆ Excellent product that I completely hate, Apr 1, 2013

By [Thirsty](#) - [See all my reviews](#)

This review is from: [Strollmaster 3000 \(Baby Product\)](#)

The Strollmaster 3000 is every parent's dream - roomy, durable, safe, and easy to fold, with a unique 17-point harness. Best yet, it weighs just 1.6 lbs. and sells for an unbelievable \$17.99. Unfortunately, it has one fatal flaw - the cupholder can only handle beverages up to 64 oz. I was dumbstruck as well. Is this America? I was left holding my 128 oz. Big Gulp like some kind of sucker. So, if you're into amazing, durable products that are a steal and virtually idyllic, then, sure, buy it. If you want to down a bathtub of Dr. Pepper, though, I'd pass.

13

Feature Engineering

- Input may be some **transformation** of the raw data
 - Because your data may not be fixed length
- 1. Because the right non-linear transformation make the problem easier
- 2. Because the right low-dimensional transform could help avoid over-fitting.
- Sometimes cleverly derived features can simplify learning.
 - E.g., House price database has length+width of room.
 - => **Linear regressor on area=L*W**, **non-linear regressor on length & width**.
- But can have **more features than original data**
- Dichotomy between designing exactly the right feature (#2) and designing very many features (#1) above.
 - If we include/make many features in the hope of finding a good one....
 - We may not know which are relevant
 - Risk of over-fitting

14

Overview

- Feature Design
- Why Reduce Dimensions?
- Feature Selection and Methods
 - Filtering
 - Wrapper
 - Built-in
- Dimensionality Reduction
 - PCA
- Summary

16

Many Dimensions

- Suppose we are given or have designed our features. Then...
- We often end up with many dimensions
 - Because we over-killed on feature engineering
 - Because it's a problem where we have very little prior knowledge, so had no choice but to include everything.
 - E.g., drug discovery, genome analysis

17

Why Reduce Dimensions?

- “Curse of Dimensionality”
- Irrelevant Data
- Computation Time
- Visualization
- Interpretation
- Many applications have $> 10^6$ features (columns)

$f(\text{spam email}) = \text{spam}$
 $f(\text{person's face}) = 32\text{yr}$

18

Why Reduce Dimensions? “Curse of Dimensionality”

- Human intuition breaks down in high dimensions
 - E.g., Gaussian, Cube, Sphere-Cube.
 - Everything is similarly far away
- So do many machine learning algorithms...
 - Slow
 - Inaccurate
 - Makes overfitting very easy, so poor generalization
- E.g.,
 - KNN: Not robust to high dimensions.
 - Linear Regression. Need data $>$ dimensions.

19

Why Reduce Dimensions? Irrelevant Data

- Curse of Dimensionality especially dangerous if
 - Many weakly relevant dimensions
 - Some very relevant, but many irrelevant dimensions
- Irrelevant dimensions, e.g.,:
 - Given article content: Classify sports versus technology.
 - Given someone's facebook profile, what product should I advertise to them?

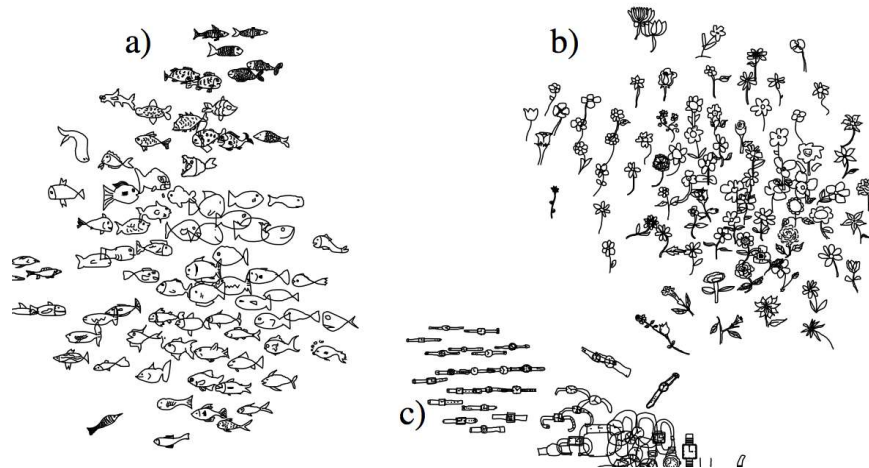
20

Why Reduce Dimensions? Computation Time

- We have seen:
 - Train: Regression $O(d^2n + d^3)$
 - Test: Regression, MaxEnt: $O(nd)$, KNN $O(nd)$.
 - (An $O(d^3)$ method will be 8x faster with $\frac{1}{2}$ the features!)
- “Big Data” / Web-scale
 - $N=10^6$, $d=10^6$
 - => Reducing dimensions is critical
- Embedded systems & mobile apps
- Real-time apps

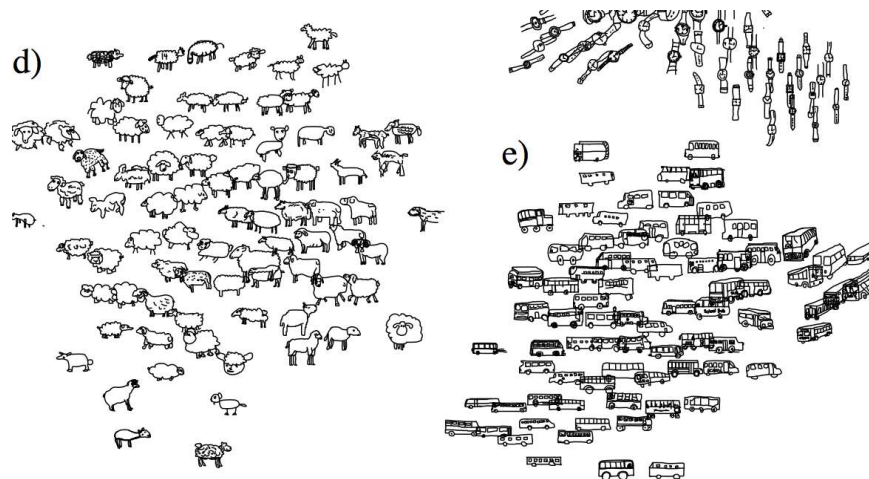
21

Why Reduce Dimensions? Visualization



22

Why Reduce Dimensions? Visualization



23

Why Reduce Dimensions? Interpretation

- Sometimes finding good dimensions is the fundamental aim
- E.g.: What causes a program to crash?
 - Features are aspects of a single program execution
 - Which branches were taken?
 - What values did functions return?
 - Classifier $F(\text{Trace})$: Crash or Not
 - Features that predict crashes well are probably bugs
- E.g.: What causes lung cancer?
 - Features are aspects of a patient's medical history
 - Binary response variable: did the patient develop lung cancer?
 - Want to legislate against features that predict lung cancer.

24

Why Reduce Dimensions? Interpretation

Task: predict chances of lung disease




Data: medical history survey

X							
Vegetarian	No						
Plays video games	Yes						
Family history	No	Reduced X	<table><tr><td>Family history</td><td>No</td></tr><tr><td>Smoker</td><td>Yes</td></tr></table>	Family history	No	Smoker	Yes
Family history	No						
Smoker	Yes						
Athletic	No						
Smoker	Yes						
Gender	Male						
Lung capacity	5.8L						
Hair color	Red						
Car	Audi						
...							
Weight	185 lbs						

25

Case Study: fMRI Brain Imaging

- “Mind Reading” experiment
- Predict mental state from voxels
 - What are they seeing/reading/thinking?
- Interesting for both prediction and insight
 - **Prediction**: ‘Mind reading’
 - E.g., Is the subject concealing information?
 - **Insight**: Which part of the brain does what
- Challenge:
 - 10-100 examples.
 - 1000k features
 - $\Rightarrow d \gg n$
 - Most dimensions irrelevant.
 - Noise (scanner, body, subject)



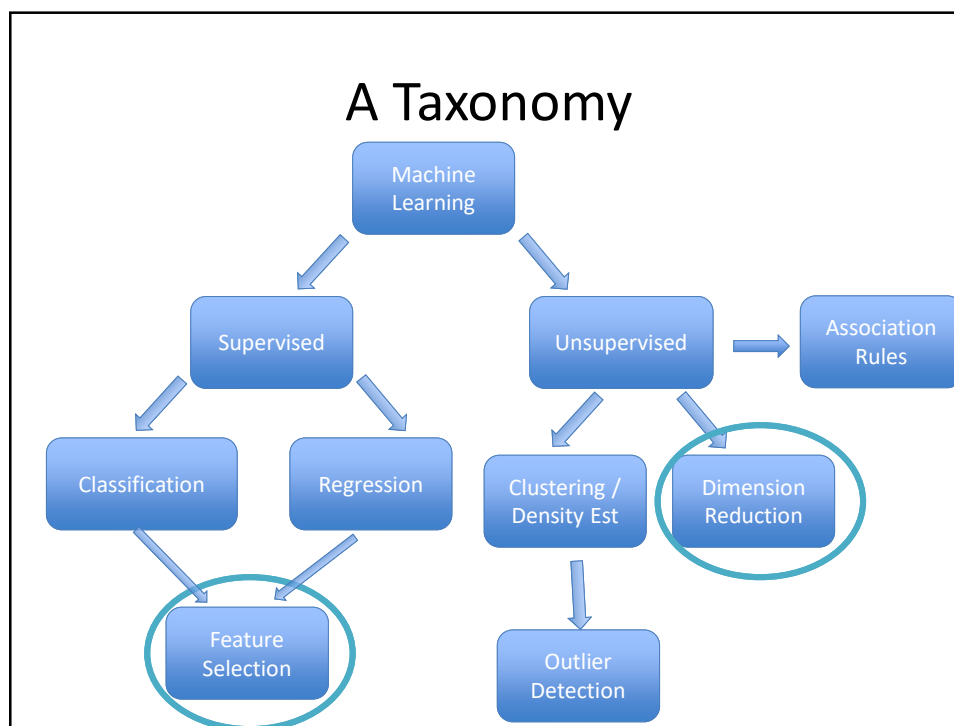
26

Reducing Dimension

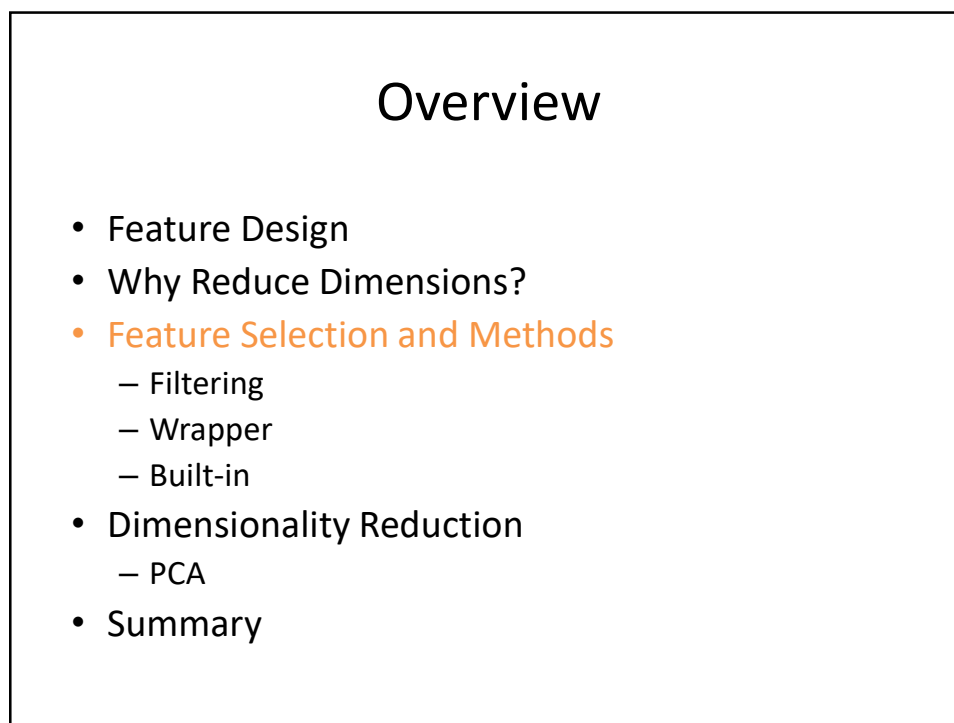
Two categories of ways to reduce dimensions....

- Feature Selection
 - Pick a good subset of features (attributes, columns), ignore the others (typically supervised)
- Dimensionality Reduction by Linear Projection
 - Transform linear combination of all features to a smaller set of features (typically unsupervised)

27



28



29

Feature Selection Methods

- Three main types
 - Filtering
 - Wrapper
 - Embedded
- Formally:
- Want to learn $y=f(\mathbf{x})$
 - $\mathbf{x}=[x_1,..x_j,..x_d]$
 - Suspect not all x_j are relevant
 - Task: Find the relevant subset
 - Challenge: there are 2^d subsets!

30

Feature Selection Methods: Filtering

- Assign a score to each feature: $s_j=\text{score}(j)$
 - Sort features j by score, and pick the top K or top $K\%$.
- Common scoring methods
 - Correlation between X_j and Y
 - Estimate the mutual information between X_j and Y

$$I(X; Y) = \sum_{y \in Y} \sum_{x \in X} p(x, y) \log \left(\frac{p(x, y)}{p(x) p(y)} \right)$$

- χ^2 test of statistical independence between X_j and Y .
- Domain Specific.
 - Text: Ignore words such as “the”, “it”.

31

Feature Selection Methods: Filtering

- Advantages:
 - Very fast
 - Simple to apply.
- Disadvantages?
 - Doesn't take into account **feature interaction**
 - => Apparently useless features can be useful when grouped together
 - It will miss these
- Practical:
 - Use light filtering as an initial step if training time is a big issue.

32

Feature Selection Methods: **Wrapper**

- Filter ignores features that can be useful in conjunction
 - Also doesn't account for limitations / power of learning algorithm.
- Wrapper methods:
 - For each subset of features:
 - Retrain learning algorithm on chosen subset
 - Evaluate learning algorithm on validation data
 - Pick the subset which has highest validation performance
- Issue:
 - Repeatedly retraining is costly
 - There are exponentially many (2^d) subsets of features.

33

Wrapper Methods: Greedy Search

Forward Selection

- Initialize no feats: $fs=\{\}$
- Do:
 - Try all unused features s
 - Find s^* to add that improves performance the most
 - Add feature s^* to fs .
- While: performance improving

34

Wrapper Methods: Greedy Search

Forward Selection

- Initialize no feats: $fs=\{\}$
- Do:
 - Try all unused features s
 - Find s^* to add that improves performance the most
 - Add feature s^* to fs .
- While: performance improving

Backward:

- Tends to find better models (interaction)
 - Frequently too expensive
- Both can be too greedy

Backward Selection

- Initialize $fs=\{1,...,d\}$
- Do:
 - Try removing each feature in fs
 - Find s^* to remove which improves performance the most
 - Remove s^* from fs .
- While: performance improving.

35

Feature Selection Methods: Embedded

- Wrapper methods:
 - Advantage: can be applied to any model (model agnostic)
 - Disadvantage: suffer from repeated re-train cost and sub-optimality (greedy).
- In some special cases, feature selection can be built into a particular learning algorithm
 - Model specific
 - May be more efficient / optimal

36

Embedded Methods

- We have seen regularization, e.g., MaxEnt and regression
 - Find $w = \operatorname{argmin} E(w, D)$
 - This is known as L2 regularization because it penalizes the squared weights

$$E_{MCLR}(w, D) = -\sum \log p(y_i | x_i) + \lambda w^T w \quad E_{MSER}(w, D) = \sum (y_i - f(x_i))^2 + \lambda w^T w$$

- Suppose some dimension j of x is totally irrelevant
 - Suppose we remove it (setting $w_j=0$)
 - No effect on the first term
 - Improves the second term
- \Rightarrow Good regularization can help with feature selection.
 - But how to achieve it?


38

Embedded Methods

- We have seen regularization, e.g., MaxEnt and regression

$$E_{MCLR}(\mathbf{w}, D) = -\sum \log p(y_i | \mathbf{x}_i) + \lambda R(\mathbf{w}) \quad E_{MSE}(\mathbf{w}, D) = \sum_i (y_i - f(\mathbf{x}_i))^2 + \lambda R(\mathbf{w})$$

– Find $\mathbf{w} = \arg\min E(\mathbf{w}, D)$



$$R_2(\mathbf{w}) = \mathbf{w}^T \mathbf{w} = \sum w_j^2 = w_1^2 + \dots + w_d^2$$

$$R_0(\mathbf{w}) = \sum I(w_j \neq 0)$$

$$R_1(\mathbf{w}) = \sum |w_j| = |w_1| + \dots + |w_d|$$

39

Embedded Methods

- We have seen regularization, e.g., MaxEnt and regression
 - Find $\mathbf{w} = \arg\min E(\mathbf{w}, D)$

$$E_{MCLR}(\mathbf{w}, D) = -\sum \log p(y_i | \mathbf{x}_i) + R(\mathbf{w}) \quad E_{MSE}(\mathbf{w}, D) = \sum_i (y_i - f(\mathbf{x}_i))^2 + R(\mathbf{w})$$

- L2 regularizer: “Ridge”
 - Fast and easy, but weak feature selection

$$R_2(w) = \sum w_j^2 = w_1^2 + \dots + w_d^2$$

- L0 regularizer is ideal
 - But very slow optimise (NP hard)
 - (because not differentiable)

$$R_0(w) = \sum I(w_j \neq 0)$$

- L1 regularizer: “Lasso”
 - Commonly chosen tradeoff.
 - Reasonably easy, reasonably quick

$$R_1(w) = \sum |w_j| = |w_1| + \dots + |w_d|$$

40


Feature Selection: When Can it Hurt?

- E.g., Fat tail in NLP
 - Many n-grams are seen only once in training.
 - 8-gram “Today I give a lecture on feature selection” only one in the mailbox, but a good predictor of WORK email.

41

Case Study: Classifying News Articles


- Approximately 10^5 English words.
 - (10^{10} bigrams, etc.)
 - Reuters article benchmark: 1500
- Rogati et al, CIKM 2002
 - Study filter methods
- Best @ 3-5% of features



42

Case Study: fMRI Brain Imaging

- Schoenmakers et al, Neurolmage 2013
 - MRI Voxel => Pixel Regression.
 - (Attempt to “mind-read” what you see from your brain activity)
 - Input: Brain voxels
 - Output: Your visual field
 - Feature selection:
 - Most brain voxels irrelevant
 - But we don’t know which...
 - Best result with L1 “Lasso” regularization



43

Case Study: fMRI Brain Imaging

- Nishimoto et al, Curent Biology, 2011

Presented clip



Clip reconstructed from brain activity



44

Case Study: Regression for Super Resolution (EECS work! 😊)

- Super Resolution: Regression problem with:
 - Input: Low resolution image
 - Output: high-resolution image
- L1 “Sparse Coding” Feature Selection



45

Case Study: View Direction Classification (EECS Work! 😊)

- Random Forest
 - Built in feat selection



46

Summary

- Sometimes we want to engineer new features:
 - Raw data may not be suitable (e.g., variable length)
 - A good derived feature may simplify the problem.
 - A suitable set of features may be lower dimensional than raw data
- Sometimes we want to select features:
 - When there are many potential inputs, and little domain knowledge to select/engineer them
 - When there are resource constraints (large scale/embedded)
 - When we engineered many features in the hope of finding a good one
 - When the feature selection is itself the goal

49

Overview

- Feature Design
- Why Reduce Dimensions?
- Feature Selection and Methods
 - Filtering
 - Wrapper
 - Built-in
- Dimensionality Reduction
 - PCA
- Summary

50


From Feature Selection to Dimensionality Reduction

- So far we selected a subset of good columns (feature selection)
 - We loose everything in the discarded columns.
- Sometimes we want to “compress” all the columns into a smaller number, but loosing the least possible information

51

Dimensionality Reduction: Overview

- Data \mathbf{x} , $|\mathbf{x}|=d$
- Derived Features \mathbf{z} , $|\mathbf{z}|=k$, $k < d$
 - $z_1 = F_1(\mathbf{x}) = x_1 + x_2$
 - $z_2 = F_2(\mathbf{x}) = 2x_3 - x_1 - x_2$
- Feature Selection
 - $z_1 = x_1$
 - $z_2 = x_3$
- Dimensionality Reduction:
 - How to find a **good linear combination** of features?
 - Restrict ourselves to **linear** combinations for now
 - Supervised: Find a linear combination that helps achieve a task.
 - Unsupervised: Find a linear combination according to some other criteria



52

Dimensionality Reduction: Linear


- Linear combinations of features can be expressed as a matrix multiply
 - $z = Ux$
- E.g., U is a binary row
 - z is a subset of x according to ones in U
- E.g., U is a list of 1s
 - z is the sum of the elements in x
- Lots more options...
 - So how to find a “good” matrix U ? Ideas?
 - Pick U that **explains** the data well / loses little information

$$\begin{matrix} 1 \\ z \end{matrix} \begin{matrix} k \\ = \end{matrix} \begin{matrix} d \\ U \end{matrix} \begin{matrix} k \\ * \end{matrix} \begin{matrix} d \\ x \end{matrix} \begin{matrix} 1 \end{matrix}$$

53

Dimensionality Reduction: Geometric Intuition

- $d=2, k=1$
- Which axis do we project to?
 - $z=x_1$?
 - $z=x_2$?




$$\begin{matrix} z \end{matrix} = \begin{matrix} U \end{matrix} * \begin{matrix} x \end{matrix}$$

54

Dimensionality Reduction: Geometric Intuition


- $d=2, k=1$
- Which axis do we project to?
 - $z=x_1$?
 - $z=x_2$?



55

Dimensionality Reduction: Geometric Intuition

- $d=2, k=1$
- Which axis do we project to?
 - $z=x_1$?
 - $z=x_2$?




56

Dimensionality Reduction: Geometric Intuition


- $d=2, k=1$
- Which axis do we project to?
 - $z=x_1$?
 - $z=x_2$?
 - New combined axis

$$z = U * x$$



57

Dimensionality Reduction (Aside: Specific Example)




58

Dimensionality Reduction: Geometric Intuition

- $d=2, k=1$
- Which axis do we project to?
 - $z=x_1$?
 - $z=x_2$?
 - New combined axis?


$$z = U * x$$



59

Dimensionality Reduction: Residual Error


- Can measure the residual error of projecting to a particular axis



60

Dimensionality Reduction: Residual Error

- Can measure the residual error of projecting to a particular axis



61

Overview

- Feature Design
- Why Reduce Dimensions?
- Feature Selection and Methods
 - Filtering
 - Wrapper
 - Built-in
- Dimensionality Reduction
 - PCA
- Examples

62


Dimensionality Reduction: Principal Components Analysis (PCA)

- PCA Objective
 - Project to the axis with **minimum residual error**
- Encoder: $\mathbf{z} = \mathbf{U}^T \mathbf{x}$
- Decoder: $\hat{\mathbf{x}} = \mathbf{U} \mathbf{z}$
- Find matrix \mathbf{U} for minimum error

$$E(\mathbf{U}) = \sum_i (\hat{x}_i - x_i)^2$$

$$= \sum_i (\mathbf{U} \mathbf{U}^T x_i - x_i)^2$$


- Same as maximize variance of \mathbf{z}



63

Dimensionality Reduction: Principal Components Analysis (PCA)


- How to solve PCA?
 - Turns out the right basis is given by **the eigenvectors of the covariance matrix**. So: $[\mathbf{U}, \mathbf{V}] = \text{eig}(\mathbf{X}\mathbf{X}^T)$
 - Faster version without explicit covariance: $[\mathbf{U}, \mathbf{S}, \mathbf{V}] = \text{svd}(\mathbf{X})$
 - Rows of \mathbf{U} are the basis
 - Diagonal of \mathbf{S} are the eigenvalues
 - Pick the first k rows
- Encode $\mathbf{z} = \mathbf{U}(1:k, :) \mathbf{x}$
- Decode $\hat{\mathbf{x}} = \mathbf{U}(1:k, :)^T \mathbf{z}$
- Note:
 - Using all the Evs will store the **exact** data



64

PCA: How to choose the number of dimensions


- How to choose?
- Each eigenvalue tells you what fraction of the variance/error is accounted for.
 - Strategy 1: Pick k dimensions.
- If you plot the eigenvalues, you typically get
 - Strategy 2: Take a number of eigenvalues such that you account for $P\%$ of the variance (E.g., $P=99\%$)



65

PCA: Pitfalls




- You must subtract the mean of the data first
 - Why?



66

PCA: Pitfalls



- You must subtract the mean of the data first
- NOT scale invariant, need to rescale first.
- 'Traditional' implementation computes $\text{Cov}(X) = XX^T$ which is already $O(nd^2)$
 - SVD can compute top k singulars in $O(ndk)$
- Second order statistics / Gaussianity assumption. This can hide many interesting patterns
- Non-linear manifolds are not covered
- Not discriminative
 - (The information your problem needs could be in a low-variance dimension)



67

PCA: Versus Linear Regression

- Regression: Predict a special output variable (y) given others (x)
- PCA: No special variable, model all the data (x) with maximum fidelity
- Error on y only versus error on all x .



68

Overview

- Feature Design
- Why Reduce Dimensions?
- Feature Selection and Methods
 - Filtering
 - Wrapper
 - Built-in
- Dimensionality Reduction
 - PCA
- Examples

69

PCA Example: Economics

- Many Economic Statistics....

Country	GDP (trillions of US\$)	Per capita GDP (thousands of intl. \$)	Human Develop- ment Index	Life expectancy	Poverty Index (Gini as percentage)	Mean household income (thousands of US\$)	...
Canada	1.577	39.17	0.908	80.7	32.6	67.293	...
China	5.878	7.54	0.687	73	46.9	10.22	...
India	1.632	3.41	0.547	64.7	36.8	0.735	...
Russia	1.48	19.84	0.755	65.5	39.9	0.72	...
Singapore	0.223	56.69	0.866	80	42.5	67.1	...
USA	14.527	46.86	0.91	78.3	40.8	84.3	...
...

70

PCA Example: Economics

- Many Economic Statistics
 - What are the underlying factors?
 - Reduce to 2 dimensions and plot...


Country	GDP (trillions of US\$)	Per capita GDP (thousands of intl. \$)	Human Develop- ment Index	Life expectancy	Poverty Index (Gini as percentage)	Mean household income (thousands of US\$)	...
Canada	1.577	39.17	0.908	80.7	32.6	67.293	...
China	5.878	7.54	0.687	73	46.9	10.22	...
India	1.632	3.41	0.547	64.7	36.8	0.735	...
Russia	1.48	19.84	0.755	65.5	39.9	0.72	...
Singapore	0.223	56.69	0.866	80	42.5	67.1	...
USA	14.527	46.86	0.91	78.3	40.8	84.3	...
...

71

PCA Example: Economics

- Many Economic Statistics
 - What are the underlying factors?
 - Reduce to 2 dimensions and plot...
 - Reveals aggregate “Wealth” + “Size” factors


Country	GDP (trillions of US\$)	Per capita GDP (thousands of intl. \$)	Human Develop- ment Index	Life expectancy	Poverty Index (Gini as percentage)	Mean household income (thousands of US\$)	...
Canada	1.577	39.17	0.908	80.7	32.6	67.293	...
China	5.878	7.54	0.687	73	46.9	10.22	...
India	1.632	3.41	0.547	64.7	36.8	0.735	...
Russia	1.48	19.84	0.755	65.5	39.9	0.72	...
Singapore	0.223	56.69	0.866	80	42.5	67.1	...
USA	14.527	46.86	0.91	78.3	40.8	84.3	...
...



72

PCA Example: Politics

- Opinion database
 - Rows: People
 - Columns: Opinions (Immigration, Crime, Tax, Welfare, Drugs, etc)
- PCA -> 1D
 - Rows: People
 - Column: Left<->Right:
 - “Lib Dem <-> Conservative”
- PCA -> 2D
 - Rows: People
 - Columns:
 - Economic & Social views



73

PCA Examples: Eigen-faces


- Each face image is a database row
 - E.g. 100x100=10000 columns.
 - What if we require to store each face image in only

$$\begin{array}{c}
 \mathbf{X}_{d \times n} \\
 \left(\begin{array}{c|c|c} \text{img}_1 & \dots & \text{img}_n \end{array} \right)
 \end{array}
 \approx
 \begin{array}{c}
 \mathbf{U}_{d \times k} \\
 \left(\begin{array}{c|c|c|c|c|c} \text{img}_1 & \text{img}_2 & \text{img}_3 & \text{img}_4 & \text{img}_5 & \text{img}_6 \end{array} \right)
 \end{array}
 \begin{array}{c}
 \mathbf{Z}_{k \times n} \\
 \left(\begin{array}{c|c|c|c|c|c} \mathbf{z}_1 & \dots & \mathbf{z}_n \end{array} \right)
 \end{array}$$

74

PCA Examples: Eigen-faces

- Each face image is a database row
- First dimensions correspond to lighting, thereafter face, hair structure
- Extensively used for face recognition
 - Speed, memory, accuracy
 - Tuck away lighting...



75

PCA Examples: Eigen-faces

- Example of reconstruction with increasing number of PCs
- Connection to general Image Compression
 - Linear vs Non-linear (eg DCT)
 - Not perceptually motivated
 - So works, but not great
- But good for revealing structure




76

PCA Examples: Other Vision..

(EECS Work ☺ - IEEE Trans KDE'11)

- Handwriting Recognition
 - Gait-based person recognition
 - Recognize identity **without subject cooperation** (CF: fingerprint / iris)
1. Feature Engineer Background Subtraction
 2. PCA to reduce dimensions



77

PCA Examples: De-anonymization

(EECS Work – Student Projects ☺)

- Recognizing People by Social Network Preferences
 - (E.g. Facebook Likes)
- Vast data matrix:
 - Rows: Persons (millions)
 - Columns:
 - Likes (binary) (billions)
- Goal:
 - Predict identity from public likes.
- Too many dimensions for most algorithms. PCA made it work.

78

PCA Examples: Text Document Classification (EECS Projects 😊)

- Category of text document?
 - Who is the **author of a text document**?
 - Computational Forensics
 - Who is the director/script-writer of a movie?
- Vast data matrix:
 - Rows: Documents
 - Columns:
 - English dictionary (100k)
- Too many dimensions for most algorithms. PCA made it work
 - (But there are much better ways to do this, more later...)

79

PCA Example: Personality

- Private traits and attributes are predictable from digital records of human behavior, Kosinski et al, PNAS 2013
- Dataset 1: (PCA!)
 - Rows: Facebook Users, Columns: Likes
- Dataset 2: (Linear regression map PCA likes => personality and demographics)
 - Rows: facebook users, Columns: Profile details & personality test.
- Outcome: **Public likes give away:**
 - Relationship status, smoking, drugs, ethnicity, voting, religion, sexual orientation, parental divorce status...
 - Intelligence, Age, Emotional Stability, Extraversion etc...

80


PCA Recap

- Plus
 - Standard tool.
 - Available in most languages / toolkits.
 - Fairly Fast and robust.
- Minus
 - Assumes linear
 - Assumes orthogonal dimensions
 - Assumes Gaussian
 - Not “discriminatively trained”

81

Beyond PCA

- Non-negative matrix factorization
 - Bases have to be positive
- Linear Discriminant Analysis & Partial Least Squares
 - Find a lower-dimension projection that **help separate classes**
- Factor Analysis
 - Don't assume the dimensions are orthogonal
- Independent Component Analysis
 - Look for the most independent basis (e.g., blind-source separation)
- Non-linear methods: Isomap, LLE, etc



82

Summary

- Feature Engineering
 - Design a small set of informative features
- Feature Selection
 - Prune irrelevant features to increase speed, reduce overfitting
 - ...and improve domain knowledge
 - E.g., Filter, Wrapper, Embedded methods
- Linear Dimensionality Reduction
 - Compress all features into a smaller number
 - E.g., PCA

83

You Should Know

- Feature Engineering
 - Some reasonable features to try in different domains
 - Especially text
- Reasons why it's useful to reduce data dimensionality
- Explain relative merits of feature selection methods
 - Filter
 - Wrapper
 - Embedded
- How feature selection applies in real-life data mining problems

84

You Should Know

- Difference between feature selection and linear dimensionality reduction.
- Linear Dimensionality Reduction
 - Explain linear dimensionality reduction as a matrix multiply
 - Encoding and decoding (compression / decompression)
- Practical
 - How to choose number of dimensions
 - Practical pitfalls to avoid
 - Some examples about how PCA can be used in real life data mining